

The background features a large, light blue watermark logo of the Technical University of Cluj-Napoca. The logo consists of a shield with a stylized 'U' and 'C' inside, with the text 'TECHNICAL UNIVERSITY' at the top, 'OF CLUJ-NAPOCA' in the middle, and 'Computer Science' at the bottom.

Fundamental Algorithms

Lecture #8
@cs.utcluj.ro

Cluj-Napoca

Computer Science

Agenda

- **Disjoint Sets**
 - Concept & list representation
 - Tree representation
- **Binomial Heaps & Binomial Trees**
 - Def
 - Basic operations
- **Fibonacci Heaps**
- **B trees**
 - Def
 - Basic operations

Disjoint Sets

- Collection of dynamic DS $S = \{S_1, \dots, S_k\}$
- \exists n elements (objects) in all k sets ($n \geq k$)
- each set S_i is identified by its *representative* element, $x \in S_i$;
- Basic operations:
 - **Build-Set** (x)
Generates a new set, with a single element \Rightarrow n sets initially, each object has its own set, and it is its own representative el.
 - **Unify** (x, y)
joins 2 disjoint sets, represented by $x \in$ and y ; builds $S_x \cup S_y$ (and destroys S_x and S_y); the representative becomes any of the 2 representatives;
 - **Find-Set** (x)
Returns a pointer to the representative element of the set containing element x .

Disjoint Sets – contd.

n = nb. of objects in the whole S

m = total nb. of operations (Build-Set, Unify, Find-Set)

$m \geq n$ (as we have n Build-Set operations)

Utility/Applications:

- speeds up execution when we need to find/group items with similar features
- graphs (connected components; MST)
- many other

Disjoint Sets - implementation

- LL
- A set = a **linked list**
- representative= the first element (head) of the list
- An object in such a list contains
 - The element from the set;
 - The pointer to the next element in the list (LL)
 - Pointer to the representative (ex: blackboard)
- **Build-Set** (builds a list with a single element) $O(1)$
- **Find-Set** (returns the representative) $O(1)$
- **Unify** (x, y) — adds x 's list at the end of y 's list;
 - representative = former y 's representative
 - all x 's elements have to update representative pointer (ex: blackboard)

Disjoint Sets – implementation – contd.

- Worst case: $O(m^2)$ for all operations
- n Build-Set (1 for each element)
- Unify
 - n times (to get to a single set)
 - $1 + 2 + 3 + \dots + 2 = O(n^2)$ (show on the blackboard)
- $n \sim m$ (actually $m > n$, yet n is linear in m)
- On average, $O(m)$ for a call of Unify, m calls, $O(m^2)$

Disjoint Sets – implementation increase efficiency

- Update pointers for the shorter lists
- Keep as knowledge their length (similar to Order Statistic Trees)
- **Theorem:** For n objects in LL with weighted unify, for m Build-Set, Find-Set and Unify takes $O(m + n \lg n)$
- **Proof:** (check the textbook – identify an **informal** justification)

Forest of Disjoint Sets

- Set = **tree** with root; keep parent pointer
- 1 node = 1 element (=1 obj) from the set
- 1 tree = a set
- The root = *representative* el.
- Basic Implem. ~ to lists (no improvement)
 - Build-Set (x) build the tree with root only
 - Find-Set (x) goes up and return the representative
 - Unify (x, y) Ex: (blackboard)

Forest of Disjoint Sets – Heuristics

(to increase performance)

- Unify based on **rank**
 - Similar to weighted unify on lists
 - The tree with less nodes will point to the tree with many nodes
 - Info kept at root level = rank = max height of the tree
 - $\text{rank} \cong \lg(\text{dim})$ (is an **approximation**, not an exact value; a guarantee that value is never exceeded)
 - Tree shrink
 - Within the **Find-Set**, each node on the *search path* will update the parent node to the *representative* (instead of parent), and leave the **rank unchanged!**
 - **Shrink does NOT change rank!** Why? Ex: blackboard
- $\text{rank} \cong \lg(\text{dim})$ It is an approximation, ONLY

Forest of Disjoint Sets – Heuristics

- $\text{Rank}[x]$
 - = max height of the subtree rooted by x
 - = nb. of edges on the longest path from x to a leaf
 - $\text{rank}[\text{leaf}] = 0$
- Find-Set leaves ranks unchanged

Forest of Disjoint Sets - Implementation

Build-Set (x)

```
p[x] <- x  
rank[x] <- 0
```

Reunion (x, y)

```
Unify (Find-Set(x), Find-Set(y))
```

Unify (x, y)

```
if rank [x] > rank [y]  
    then p[y] <- x  
    else p[x] <- y  
if rank [x] = rank [y]  
    then rank [y] = rank [y] + 1
```

Forest of Disjoint Sets - Implementation

Find_Set(x)

```
if x != p[x]  
  then p[x] <- Find_Set(p[x])  
return p[x]
```

Binomial Trees

- Degree-based augmented trees; denoted B_k
- Degree of the tree = number of descendants of the tree
- Properties of a B_k tree:
 - P1: Degree of the root (B_k) = k ;
 - P2: Number of nodes (B_k) = 2^k ;
 - P3: Height(B_k) = k ;
 - P4: Number of nodes at level i in (B_k) is $C_k^i = \binom{k}{i}$
 - P5: If children of (B_k) are numbered from the left to the right as $(k-1, k-2, \dots, 0)$, then child i is a (B_i) tree
- Recursive definitions:
 - A B_k tree is 2 B_{k-1} - trees, with their roots linked (picture on the blackboard)
- Be aware of the implication the equivalence definitions following P5 and the recursive definition!

Binomial Trees

- **Recursive definitions:**
 - A B_k tree is 2 B_{k-1} - trees, with their roots linked (see picture and follow discussion)
 - A B_k tree is collection of k trees: $B_{k-1}, B_{k-2}, \dots, B_1, B_0$ – trees (see picture and follow discussion)
- **Proof of P4: Number of nodes at level i of (B_k) is C_k^i**
 - Goes by induction
 - On level i we have nodes from 2 B_{k-1} trees
 - From the first tree (containing the root of B_k) #nodes at level $i = C_{k-1}^i$
 - From the second one #nodes at level $i-1$ (one level less; level measured from the root) = C_{k-1}^{i-1}
 - So there are $C_{k-1}^i + C_{k-1}^{i-1} = C_k^i$ nodes at level i in B_k

Binomial Heaps

- **Binomial Heap (H) = A set of Binomial trees with the following properties:**
 - **P1:** each node has a key;
 - **P2:** each binomial tree in H is heap-ordered (min on top);
 - **P3:** for any k, there is at most one B_k tree in H.
- **Consequence:** if H has n nodes, it has at most $\lfloor \lg n \rfloor + 1$ binomial trees.
 - **Justification:**
 - Max number of trees a H may have = one of each type
 - If each type of tree is present, the number of nodes for $B_{k-1}, B_{k-2}, \dots, B_1, B_0$ is $2^{k-1}, 2^{k-2}, \dots, 2^0$ respectively
 - Nb of nodes of H is their sum $= 2^{k-1} + 2^{k-2} + \dots + 2^0 = 2^k - 1$
 - Denote $2^k = n$. H has n nodes, and k trees ($\lfloor \lg n \rfloor + 1$)

Binomial Heaps – Operations

($|H|=n$) (for all, examples on the blackboard)

- **Make-Heap** **$O(1)$**
 - Builds an empty Binomial Heap
- **Binomial-Heap-findMinimum(H)** **$O(\lg n)$**
 - Returns the pointer to the root of the B with the min key (NOT removed);
- **Binomial-Heap-Unite(H1, H2) = merge + links** **$O(\lg n)$**
 - merge – merges 2 rooted lists (H1, H2) into a single one sorted by degree (increasing order) **$O(\lg n)$**
 - link – changes a pair of B_{k-1} trees into a B_k tree **$O(1)$**
 - Unite = merge + links from left to right **$O(\lg n) + \lg n O(1)$**

Binomial Heaps – Operations


($|H|=n$) (for all, examples on the blackboard)

- **Binomial-Heap-extractMinimum(H)** **$O(\lg n)$**
 - **Binomial-Heap-findMinimum(H)** **$O(\lg n)$**
 - **removes that tree from H** **$O(1)$**
 - **Make a heap out of the binomial tree containing the min key $\Rightarrow H1$** **$O(\lg n)$**
 - **Binomial-Heap-Unite(H, H1)** **$O(\lg n)$**
- **Binomial-Heap-keyDelete(H1, x)** **$O(\lg n)$**
 - **As if we were to extract min.**
 - **How?**
 - **Decrease the key to delete to $-\infty$** **$O(1)$**
 - **Restore the heap property of the Binomial tree +** **$O(\lg n)$**
 - **Extract min** **$O(\lg n)$**

Computer Science

Fibonacci Heaps

- **Collection of ordered trees**
- **Properties:** like Binomial Heaps with some constraints added/removed.
- **Relaxed constraints:**
 - May contain several trees of the same degree
 - Rooted, yet unordered
- **Added constraints:**
 - Children at a given level in a tree are linked to each other (left/right) in a circular, doubly linked list (child list)
 - Node augmentation: $\text{degree}[x] = \# \text{ of children in the child list of } x$
 - The heap maintains a pointer ($\text{min}[H]$) to the root of the tree containing the min key.



Fibonacci Heaps

- **Collection of ordered trees**
- **Properties:** like Binomial Heaps with some constraints added/removed.
- **Relaxed constraints:**
 - May contain several trees of the same degree
 - Rooted, yet unordered
- **Added constraints:**
 - Children at a given level in a tree are linked to each other (left/right) in a circular, doubly linked list (child list)
 - Node augmentation: $\text{degree}[x] = \# \text{ of children in the child list of } x$
 - The heap maintains a pointer ($\text{min}[H]$) to the root of the tree containing the min key.
- **Operations:**
 - Like for Binomial Heaps (Homework!)
 - Obs: due to relaxations, operations faster:
 - Insert – a node which is a tree, at the top level
 - FindMin – has a pointer
 - Union – simpler, since they are not ordered
 - DecreaseKey – $O(1)$, as opposed to $O(\log n)$

Binomial/Fibonacci Heaps

Comparative analysis

Operation	Binomial	Fibonacci
Make heap	$O(1)$	$O(1)$
Insert	$O(\lg n)$	$O(1)$
Find min	$O(\lg n)$	$O(1)$
Extract min	$O(\lg n)$	$O(\lg n)$
Union	$O(\lg n)$	$O(1)$
Decrease key	$O(\lg n)$	$O(1)$
Delete	$O(\lg n)$	$O(\lg n)$

Computer Science

B-trees

- Previous DS reside in the primary memory
- **Trees on secondary storage devices (disk)**
- A node may have many children
- Goal: **decrease the number of pages accessed** when search for a node
- Store a very large number of keys
- Maintain the height of the tree under control (h very small)

B-trees – contd.

- **Typical pattern while working with B-trees:**

`x` ← pointer to some object

Disk-Read(`x`)

Operations that access/modify some fields of `x`

Disk-Write(`x`)

- **Once in memory, operations are performed fast**
- **Objective: as few pages read/write operations**

Computer Science

B-trees – contd.

Generalization of BST (with ordered lists)

- P1: $n[x]$ keys in node x
- P2: keys are ordered

$key_1[x] \leq key_2[x] \leq \dots \leq key_{n[x]}[x]$

- P3: An internal node contains $n[x]+1$ pointers to the children $c_i[x]$
- P4: is a search tree:

$key[c_1[x]] \leq key_1[x] \leq key[c_2[x]] \leq \dots$

- P5: All leaves are at the same level = height of the tree = h
- P6: t = degree of the tree; $\min(t)=2$. Every node (except for the root) has at **least $t-1$ keys**, and **t children**
- P7: Every node (except for the root) has at **most $2t-1$ keys**, and **$2t$ children**

B-trees – Search

- **One pass procedures (top-down ONLY, NO back up; as opposed to PBT, AVL, RB trees)**
- **For ALL operations, the process is JUST top->down!!!**
- **Search**
 - Straightforward generalization of BST search
 - Combined with ordered list search
 - #pages accessed (worst) $O(h)$
 - Access time in a page (worst) $O(t)$
 - **Overall $O(th)$**

B-trees – Insert

- **One pass procedures (top-down ONLY, NO back up updates; as opposed to PBT, AVL, RB trees)**
- **Insert**
 - **Search** for a LEAF position to insert
 - Insertion is performed in an **EXISTING** leaf
 - Along the path while searching (top-down), ensure **there is space for a safe insert** (*split full nodes* on the path down to avoid overflowing, so that the insertion is successful in an existing node!!!)
 - A key added to a **full node** (leaf included) will induce:
 - the **migration of the median key to the parent node**,
 - and the **split** of the given node into 2 nodes (leaf into 2 leaves)
 - **Time: $O(h)$**
 - **$O(h)$ disk accesses,**
 - **$O(t)$ CPU time in one page**

B-trees – Insert

- Like in any BST tree insert in a leaf (in a leaf, NOT as leaf; the node is NOT now created!)
- Stages:
 - **Search** the path for the position (leaf) to insert
 - Ensure the search **path is safe**
 - **Insert** the key in the corresponding **leaf**
- Types of nodes to **store** a key:
 - Leaf (key to be inserted)
 - Non-leaf (the “safe path” step = in the attempt to make room = in the split stage with median migration up)
- Possible issues
 - Attempt to store in a full node, with $(2t-1)$ keys – issue – **node overflows! Not allowed.**
- Cases to analyze
 - Not overflowing node – no issue
 - Overflowing node – issue – needs a strategy to handle it

B-trees – Insert

- **Strategy**

- Along the searched path, ANY full node along the path (with already $(2t-1)$ keys) is “fixed” (allowing for a potential full node to accept a new key to be added):
 - **Divide the full node** in 2 nodes with $(t-1)$ keys
 - The **median key** in the full node is **promoted to the parent** node (there *is room*, as we proceed top-down, and an upper node was “fixed”, is not overflowing)
 - if ***root is full***, **increase the height** (by adding 1 more node = new root); the ONLY case of *height increase*.

- **Insert procedure**

- Top-down approach (descendent)
- There is **NO** operation performed on return (bottom up)

B-trees – Delete

- **One pass procedures (top-down**, NO back up update; as opposed to AVL, RB trees)
- **Delete**
 - **Search** for the node containing the key to be removed and identify its type (leaf/no leaf – all nodes are either internal, or leaves; the tree is complete) (z from BST)
 - Physical ***removal*** of one object which *belongs to a leaf* (y in BST)
 - Q: Why y, a node with one only child in a BST is in a leaf in a B-tree?
 - Along the path while searching (top-down), **ensure the constraints for a safe delete are met** (*merge nodes with degree t on the path down to avoid underflowing, so that the deletion is possible*)
 - **Time: $O(th)$**
 - $O(h)$ disk accesses
 - $O(t)$ CPU time in one page =>

B-Trees – delete

- **Like in ANY other BST tree**
 - **Search for the key to remove** (pointed by **z**)
 - **If in leaf, delete it**
 - **If not in the leaf, remove (physically) a node** (pointed by **y**; its content is moved in the node pointed by **z**) **with one-single child** (pointed by **x**) - **the predecessor/successor** – (in B-trees **y** is in a leaf, hence **x** points to nil)
- **Types of nodes containing the key to be deleted**
 - Leaf
 - Non-leaf (i.e. internal)
- **Possible issues**

Node type	Capacity
Leaf	Does not underflow
Non-leaf	Underflows

 - **Attempt to delete from a node with only $(t-1)$ keys – issue – node underflows! Not allowed**
- **Cases to analyze**

B-Trees – delete

- **Cases to analyze**

Issue: attempt to delete from a node with only $(t-1)$ keys – node underflows! Not allowed

- **Solution**

- Similar strategy as in case of insert: **prevent, rather than repair**
- In the **search stage (for the key to delete)**, ensures that **each** node (along the searched path) has the ability to allow for delete (does not underflow)
- **Along the searched path, any node with only $(t-1)$ keys is “fixed”**
 - If any of the **sibling** nodes has **at least t keys**, “**borrow**” a key from it (promote the last/first key from the left/right brother node to the parent node, and move down the appropriate key from parent to the almost underflowing node). (see examples for case 3.1 – next slide)
 - If **both sibling** neighbors have **only $(t-1)$ keys**, **merge** the underflowing node with one of the 2 siblings, and **put the key from the parent node in between them in the new generated (by merge) node**. (see examples for case 3.2.1 next slide). Maybe a height shrink occurs (see examples for case 3.2.2 next slide).
- **Delete procedure**
 - Top-down approach (descendent)
 - There is no operation performed on return (bottom up)

Node type	Capacity
Leaf	Does not underflow
Non-leaf	Underflows

Node type	Capacity
Leaf	Does not underflow
Non-leaf	Underflows

B-Trees – delete

- Cases – only 3 to be considered (not 4): non-leaf/underflow is not considered (since along the path, the underflow situation is solved, anyway).
- Cases: (follow the examples – blackboard)
 - Case1 – key in leaf, node does not underflow
 - Simply delete the key
 - Case2 – key not in leaf, node does not underflow
 - Remove pred/succ (from a leaf for a BT) like in BST. The case reduces to case 1 or 3.
 - Case3 – key in leaf, node underflows
 - 3.1 sibling consistent (at least one sibling neighbor has at least t keys) sol: “borrow” from sibling
 - Promote the last/first key in consistent neighbor to parent
 - Move down the key in parent node in between the leaf and consistent sibling to the underflowing leaf
 - 3.2 neighbor siblings both with just $(t-1)$ keys
 - Merge the underflowing leaf with one sibling neighbor by adding in the middle the key in the parent in between the leaf and sibling
 - 3.2.1 keep the height of the tree
 - 3.2.2 *decrease the height* of the tree (if the *parent* is the *root* and has just one single key)